Operator’s Manual

Access PoD - Volume 2 - Public Review Draft

Document identifier: APOD-TR-003 — Version B — Artifact Pack v1b
Document Status: Informative
Audience: Operators, Governance Leads, Compliance Leads, Risk Owners
Scope: Governance operation and regulator engagement
Out of Scope: Cryptographic mechanisms, system internals, onboarding code

---

1. Purpose of This Manual

This Operator’s Manual explains how to **use the Version B artifact pack** to:

- operate runtime AI governance using Access PoD,
- prepare for regulatory or standards engagement (including EU AI Act Annex VII),
- generate, maintain, and interpret governance evidence, and
- communicate assurance outcomes without asserting certification or compliance.

This manual does **not** teach you how to build or deploy AI systems.
It teaches you how to **govern them responsibly and demonstrably**.

---

2. What Version B Is (and Is Not)

2.1 What Version B Is

Version B is an **architecture-specific governance instantiation** that:

- demonstrates how governance controls attach to execution boundaries,
- shows how evidence is produced at runtime,
- enables continuous, revocable assurance via CSC,
- supports regulator-ready evidence exchange.

Version B uses **Q-Pathformer as a reference implementation**, but the governance model is **architecture-agnostic**.

2.2 What Version B Is Not

Version B is **not**:

- a certification scheme,
- a conformity assessment,
- a legal compliance determination,
- a security or cryptography specification,
- an onboarding or deployment guide.

Any claim of “certified”, “approved”, or “compliant” based on these artifacts alone is **invalid**.

---

3. Artifact Pack Overview and Regulatory Use

This section explains the structure of the Version B artifact pack, the role of each artifact class, and how operators should use them during regulatory or standards engagement.
The artifact pack is intentionally structured to separate **authoritative machine-readable governance evidence** from **human-readable guidance**, ensuring clarity, auditability, and scope discipline.

3.1 Artifact Classes Overview

The Version B artifact pack contains **three distinct classes of artifacts**, each serving a different purpose.

3.1.1 Native Governance Artifacts (Machine-First)

These are the **authoritative runtime governance artifacts**.
They represent the system’s declared governance state and evidence availability.
These artifacts are:

- machine-readable,
- schema-validated,
- scope-bound and time-bound, and
- suitable for ingestion by Access PoD or regulator tooling.

Files in this class include:

- apod.self_assessment.eu_annex_vii.v1b.json
- apod.self_assessment.eu_annex_vii.schema.json
- apod.evidence_manifest.json
- apod.evidence_manifest.schema.json
- csc_snapshot.json
- csc_snapshot.schema.json

These files are the **single source of truth** for governance status and assurance signals.

3.1.2 Expanded / Human-Readable Evidence Aids

These artifacts support **operator preparation and regulator understanding**, but they are **not authoritative governance records**.
They are used to:

- guide internal teams,
- explain governance intent,
- align discussions with regulatory language, and
- prepare structured responses.

Files in this class include:

- EU AI Act Annex VII-Aligned Evidence Checklist (PDF/DOCX)
- Operator Evidence Checklist – Pre-Regulatory Engagement (PDF/DOCX)
- Expanded / non-native self-assessment JSON (review matrix)

These artifacts **support interpretation** but do not replace native JSON files.

3.1.3 Operator Guidance and Process Documentation

This class provides **procedural clarity** for operators.
It explains:

- how to use the artifact pack,
- how to engage regulators appropriately, and
- how to avoid misrepresentation or scope creep.

Files in this class include:

- Operator's Manual (this document)
- Regulator Briefing (Informative)
- Regulator Q&A Playbook
- Certification Differentiation Insert

These artifacts **do not constitute governance evidence** and must not be presented as such.

3.2 Authoritative vs Supporting Artifacts

Operators must clearly distinguish between:

Category	Purpose	Authority
Native JSON artifacts	Governance state & assurance	**Authoritative**
Evidence checklists	Readiness & alignment	Supporting
Manuals & briefings	Process & interpretation	Informative

If a discrepancy exists, **native JSON artifacts take precedence**.

3.3 Regulator Evidence Request Response Template

(Pre-Approved Wording)
This template provides **approved, regulator-safe language** for responding to written evidence requests.
Operators are encouraged to reuse this wording verbatim where appropriate.

-See Section 13 for the pre-approved response template-

3.4 Operator Guidance: What to Do When Evidence Is Requested

When responding to regulators:

- **Always provide native JSON artifacts first**
- Use checklists and summaries to explain, not replace, evidence
- Never back-fill evidence after the fact
- Never claim “certified”, “approved”, or “compliant”
- Treat degradation or revocation as normal governance signals

If evidence does not exist, state so clearly.

Section 3 Summary

- The artifact pack is intentionally divided into three classes.
- Native JSON artifacts are authoritative.
- Human-readable documents are supportive, not evidentiary.
- Pre-approved response wording protects against over-claiming.
- Clear boundaries preserve trust and regulatory credibility.

---

4. Operator Roles and Responsibilities

Every operator should clearly assign the following roles:

Role	Responsibility
Governance Lead	Owns Access PoD configuration and policy decisions
Risk Owner	Oversees SAPP simulations and risk categorisation
Assurance Lead	Oversees CSC issuance, degradation, revocation

Role	Responsibility
Compliance Lead	Interfaces with regulators and standards bodies
Engineering Lead	Ensures execution APIs emit required telemetry

Responsibility does not transfer to vendors, models, or APIs.

---

5. Operating the Governance Lifecycle

Step 1 — Define Scope and Intended Use

Before any regulator engagement:

- confirm system ID, environment, jurisdiction, and intended use,
- confirm whether *high-risk classification is claimed* (EU AI Act),
- record scope in apod.self_assessment.eu_annex_vii.json.

If scope is unclear, **stop here**.

Step 2 — Complete the Self-Assessment (Native)

Populate:

apod.self_assessment.eu_annex_vii.v1b.json

Rules:

- Only mark **“met”** if evidence exists.
- Use **“partially_met”** liberally — regulators prefer honesty.
- “not_applicable” is acceptable where justified.

This file is your **operator declaration**, not a compliance claim.

Step 3 — Index Evidence (Manifest)

Populate:

apod.evidence_manifest.json

Rules:

- List **what evidence exists**, not the evidence itself.
- Use summaries, IDs, timestamps, scopes.
- Never embed restricted content.

If evidence is not indexed, **it does not exist for governance purposes**.

Step 4 — Generate CSC Snapshot

Generate:

csc_snapshot.json

This file:

- is computed from telemetry,
- is time-bound and scope-bound,
- may degrade or revoke automatically.

Rules:

- CSC stars ≠ approval.
- CSC can go down. That is expected.
- Never cache or reuse snapshots outside their validity window.

Step 5 — Internal Readiness Check

Before regulator contact, complete:

- **Operator Evidence Checklist**
- **EU Annex VII-Aligned Checklist**

If more than ~30% is "not_met", **do not engage yet**.

---

6. Preparing for Regulator Engagement

What You Share

You may share:

- native JSON artifacts,
- checklist PDFs,
- CSC snapshots,
- scope-bound summaries.

What You Do NOT Share

You must not share:

- model weights,
- training code,
- cryptographic keys,
- validation node internals,
- thresholds or playbooks.

If asked, state:

“This framework provides evidence of governance behavior, not system internals.”

This is correct and defensible.

---

7. Interpreting CSC Results

CSC represents **current governance risk posture**.

Star Change	Meaning
Upgrade	Controls improved, evidence strengthened
Degradation	Telemetry gaps, drift, or control weakness
Revocation	Trust unsupported — protective action

A revocation is not failure.
Failure is hiding it.

---

8. Handling Degradation and Revocation

When CSC degrades or revokes:

1. Do **not** override manually.
2. Identify evidence gaps or failures.
3. Run SAPP simulations if applicable.
4. Restore controls and evidence.
5. Issue a **new CSC snapshot**.

Never “re-issue” without new evidence.

---

9. Relationship to Version C (Onboarding)

This manual governs **Version B only**.

Version C will cover:

- Q-Pathformer onboarding,
- operational workflows,
- execution mechanics.

Do not mix Version C content into regulator discussions unless explicitly requested.

---

10. Common Operator Mistakes (Avoid These)

- Treating CSC as certification
- Back-filling evidence after incidents
- Over-marking “met”
- Sharing restricted internals
- Letting engineering override governance
- Treating governance as paperwork

---

11. Final Operator Principles

- **If telemetry does not exist, assurance does not exist.**
- **Governance must operate at runtime, not review time.**
- **Trust is earned continuously, not granted once.**

- **Evidence beats explanation.**

---

12. Operator Sign-Off Checklist

Before engagement, confirm:

- Scope is defined and frozen
- Self-assessment completed honestly
- Evidence manifest is complete
- CSC snapshot is current
- Teams understand boundaries
- No certification claims are made

---

13. Regulator Evidence Request Response Template

(Informative — Pre-Approved Wording)

This one-page template is intended for **written responses** to regulators, notified bodies, or standards organizations requesting governance evidence.

Operators are encouraged to reuse this wording verbatim.

Subject: Response to Governance Evidence Request — Access PoD Framework (Version B)

Thank you for your request regarding governance and assurance evidence for the referenced AI system.

In response, we provide the following materials and clarifications:

---

1. Governance Evidence Provided

We provide governance evidence in the form of:

- machine-readable governance self-assessment records,
- a scope-bound evidence manifest indexing available artifacts, and
- a time-bound Compliance Star Certificate (CSC) snapshot derived from runtime telemetry.

These artifacts reflect **observed governance behavior**, not declared intent or static documentation.

---

2. Nature of the Compliance Star Certificate (CSC)

The Compliance Star Certificate (CSC):

- is a **continuous assurance signal**, not a certification or approval,
- represents current governance risk posture within the declared scope and validity window, and
- may degrade or revoke automatically when evidence no longer supports prior trust claims.

CSC outputs must not be interpreted as regulatory authorization or compliance status.

---

3. Scope and Boundaries of Disclosure

The evidence supplied:

- is limited to governance execution, provenance, telemetry, and assurance signals,
- does not include model weights, training code, or proprietary internals, and
- does not expose cryptographic mechanisms, validation node internals, or security playbooks, which are governed separately.

Governance responsibility remains with the system operator.

---

4. Human Oversight and Accountability

Human-in-the-loop oversight, escalation pathways, and governance decision records are included where applicable.

Delegation to external models or services does not transfer responsibility for governance outcomes.

---

5. Limitations

The materials provided do **not** constitute:

- regulatory certification,
- conformity assessment,
- legal compliance determination, or
- authorization for deployment.

They are supplied to support **informed regulatory review of runtime governance effectiveness**.

Should further clarification be required, we are available to provide **additional scope-bound summaries or explanations**, subject to confidentiality and governance constraints.

Sincerely,
[Name]
[Role — Governance / Compliance Lead]
[Organization]
[Contact Information]

Sincerely,
[Name]
[Role — Governance / Compliance Lead]
[Organization]

---

14. Executive Do / Don't Guide for Regulatory Engagement

(Informative — Executive Reference)

This page is intended for **executives, board members, and senior leaders** who may speak with regulators, standards bodies, or public authorities without deep technical context.

Its purpose is to **prevent misstatements**, over-claims, or scope confusion.

DO — What Executives Should Say and Emphasize

DO emphasize evidence and behavior

- "We provide runtime governance evidence, not static declarations."
- "Assurance is derived from observed system behavior."

DO explain CSC correctly

- "The Compliance Star Certificate is a continuous assurance signal."
- "It reflects current governance risk posture, not approval or certification."

DO acknowledge degradation openly

- "Assurance can degrade or revoke if evidence weakens."
- "That is a protective feature, not a failure."

DO state boundaries clearly

- "We do not provide model internals or proprietary code."
- "Governance responsibility remains with the operator."

DO align with regulatory intent

- "This framework complements existing regulation."
- "It supports ongoing oversight rather than one-time approval."

DON'T — What Executives Must Avoid Saying

DON'T claim certification or compliance

- ❌ "We are certified under this framework."
- ❌ "This proves we comply with the AI Act."

DON'T equate CSC with approval

- ❌ "We have a four-star rating, so the system is approved."
- ❌ "Higher stars mean the model is safer or better."

DON'T promise permanence

- ❌ "Our assurance level is stable."

- ❌ “Once achieved, trust remains valid.”

DON’T disclose restricted material

- ❌ Model weights or training data
- ❌ Cryptographic mechanisms or keys
- ❌ Incident thresholds or response playbooks

DON’T shift responsibility

- ❌ “The model/vendor is responsible.”
- ❌ “The architecture ensures compliance.”

Executive Anchor Statement (Use This When Unsure)

**“We do not ask regulators to trust claims or documentation.
We provide evidence of how governance operates in practice, and that evidence can change over time.”**

This statement is:

- accurate,
- defensible,
- regulator-aligned, and
- safe across jurisdictions.

End of Operator’s Manual