

Regulator Q&A Playbook

(Anticipated Questions + Approved Answers)

Document identifier: APOD-TR-003 — Version B — Artifact Pack v1b

This playbook is designed for:

- regulator meetings,
- written follow-ups,
- standards consultations, and
- conformity-adjacent discussions (EU AI Act, ISO, sector regulators).

All answers are **carefully bounded** — no over-claims, no leakage, no certification language.

Q1. Is this a certification scheme or conformity assessment?

Approved answer:

No. This framework is **not a certification scheme** and does not issue approvals or compliance status.

It provides **runtime governance infrastructure** and **evidence-derived assurance signals** that regulators may observe and interpret alongside existing regulatory tools.

Q2. What does the Compliance Star Certificate (CSC) represent?

Approved answer:

CSC is a **computed, time-bound assurance signal** derived from observed governance telemetry.

It represents **current governance risk posture**, not system quality, performance, or regulatory approval.

CSC can degrade or revoke automatically when evidence degrades.

Q3. *Why should regulators trust CSC if it's not a certification?*

Approved answer:

CSC is not trusted because it is declared — it is trusted **only when supported by evidence**.

If telemetry does not exist, CSC does not exist. If evidence weakens, CSC degrades. This makes it harder to game than static certification.

Q4. *Do you expect regulators to review model weights or training code?*

Approved answer:

No. This framework is intentionally designed so regulators **do not need access to model internals**.

Oversight focuses on **behavior, control effectiveness, and accountability**, not proprietary implementation detail.

Q5. *How does this align with EU AI Act Annex VII requirements?*

Approved answer:

The framework aligns by making **risk management, post-market monitoring, human oversight, and traceability observable at runtime**.

It complements Annex VII by replacing static documentation with **continuous evidence**, without redefining legal obligations.

Q6. *Who is accountable if something goes wrong — the model, the architecture, or the operator?*

Approved answer:

Accountability always remains with the **operator asserting governance conformance**.

Delegation to models, APIs, or architectures does not transfer responsibility.

Q7. What happens if assurance degrades or is revoked?

Approved answer:

Degradation or revocation is a **protective control**, not a failure.

It signals that trust claims are no longer supported by evidence and that corrective action is required before assurance can be re-issued.

Q8. Can operators override or manually adjust CSC scores?

Approved answer:

No. CSC is **computed from telemetry**. It cannot be manually overridden or administratively reinstated without new evidence.

Q9. How do you prevent systems from hiding problems to keep a high score?

Approved answer:

The framework enforces **non-compensatory scoring** and **telemetry-dependency**.

Missing telemetry constrains or suspends assurance. Hiding data lowers trust rather than preserving it.

Q10. Why use Q-Pathformer as a reference implementation?

Approved answer:

Q-Pathformer is used **only as a reference execution substrate** to demonstrate feasibility.

Governance authority does not reside in the architecture, and adoption of Q-Pathformer is not required to use the framework.

Q11. Are cryptographic mechanisms used? Can regulators inspect them?

Approved answer:

Cryptographic mechanisms support integrity and provenance but are governed by **restricted specifications**.

Regulators are not required to inspect cryptography to evaluate governance effectiveness.

Q12. *How is catastrophic data or model overwrite prevented?*

Approved answer:

Execution is constrained through **policy-bound APIs** and governance-authorized change workflows.

No component can propose, approve, and execute changes unilaterally.

Q13. *Can this framework be used across vendors and sectors?*

Approved answer:

Yes. The framework is **architecture-agnostic and sector-aware**.

It standardizes governance evidence, not system design.

Q14. *What evidence can regulators reasonably request?*

Approved answer:

Regulators may request:

- runtime governance telemetry summaries,
- simulation (SAPP) risk summaries,
- provenance and lineage semantics,
- CSC snapshots with scope and timestamps, and
- records of degradation or revocation.

Regulators should not expect model internals, cryptographic keys, or proprietary datasets.

Q15. *What happens next if this framework is adopted?*

Approved answer:

The framework is intended to evolve through **validated practice and standards engagement**, not unilateral declaration.

Version C will address onboarding and operational detail without redefining governance authority.

Final guidance for operators

When in doubt, anchor on this sentence:

“This framework enables regulators to observe how governance operates in practice, not how models are engineered internally.”

That sentence is:

- accurate,
- defensible,
- standards-aligned, and
- repeatable across jurisdictions.